

MOTIVATION

In the past, I have written a little library, I call `mato` - short for Markdown Transformer framework¹. In that library, I make heavy use of `groff` - GNU roff - which does the job of transforming output, my library generates, into PDF files. It does that fairly quickly, say around half a second for the average document.



But this step is by far the most expensive step in a whole chain of transformation steps: first `mato` reads the input file, and creates a abstract syntax tree from it. That tree is being processed several times and then rendered into `groff` input format. After that `groff` takes over. It is itself built in a similar way as a compiler, it has a front-end doing the parsing and a backend, doing the output generation. These backends are called *devices* in `groff` parlance.

One such device, I use heavily is the PDF device, implemented by `gropdf.pl`. As the name suggests, it is written in Perl² and thus rather expensive with regards to CPU consumption.

Because, I wanted to build something in zig, I decided, this would be a good toy project, to practise my skills.

FIRST ITERATION - GETTING SOMETHING

My first goal is to see some output in macOS' preview app. I have a sample PDF file copied over from a book called "PDF Explained" by O'Reilly. That book served as a starting

¹I now notice for the first time, that the `o` character is out of place. It is a historical left-over, because initially, I called it `matote`, Markdown to TeX converter, but at some point dropped the TeX and switch it to `groff`, because it is much simpler and much quicker.

²Some of you might know this scripting language still!

point, but I soon learned, that the Adobe “PDF Reference 1.7” (see [here](#)).

The first problem my current implementation has, is that the output stream is chosen in the pdf implementation. So, the first challenge becomes, factoring that out.